

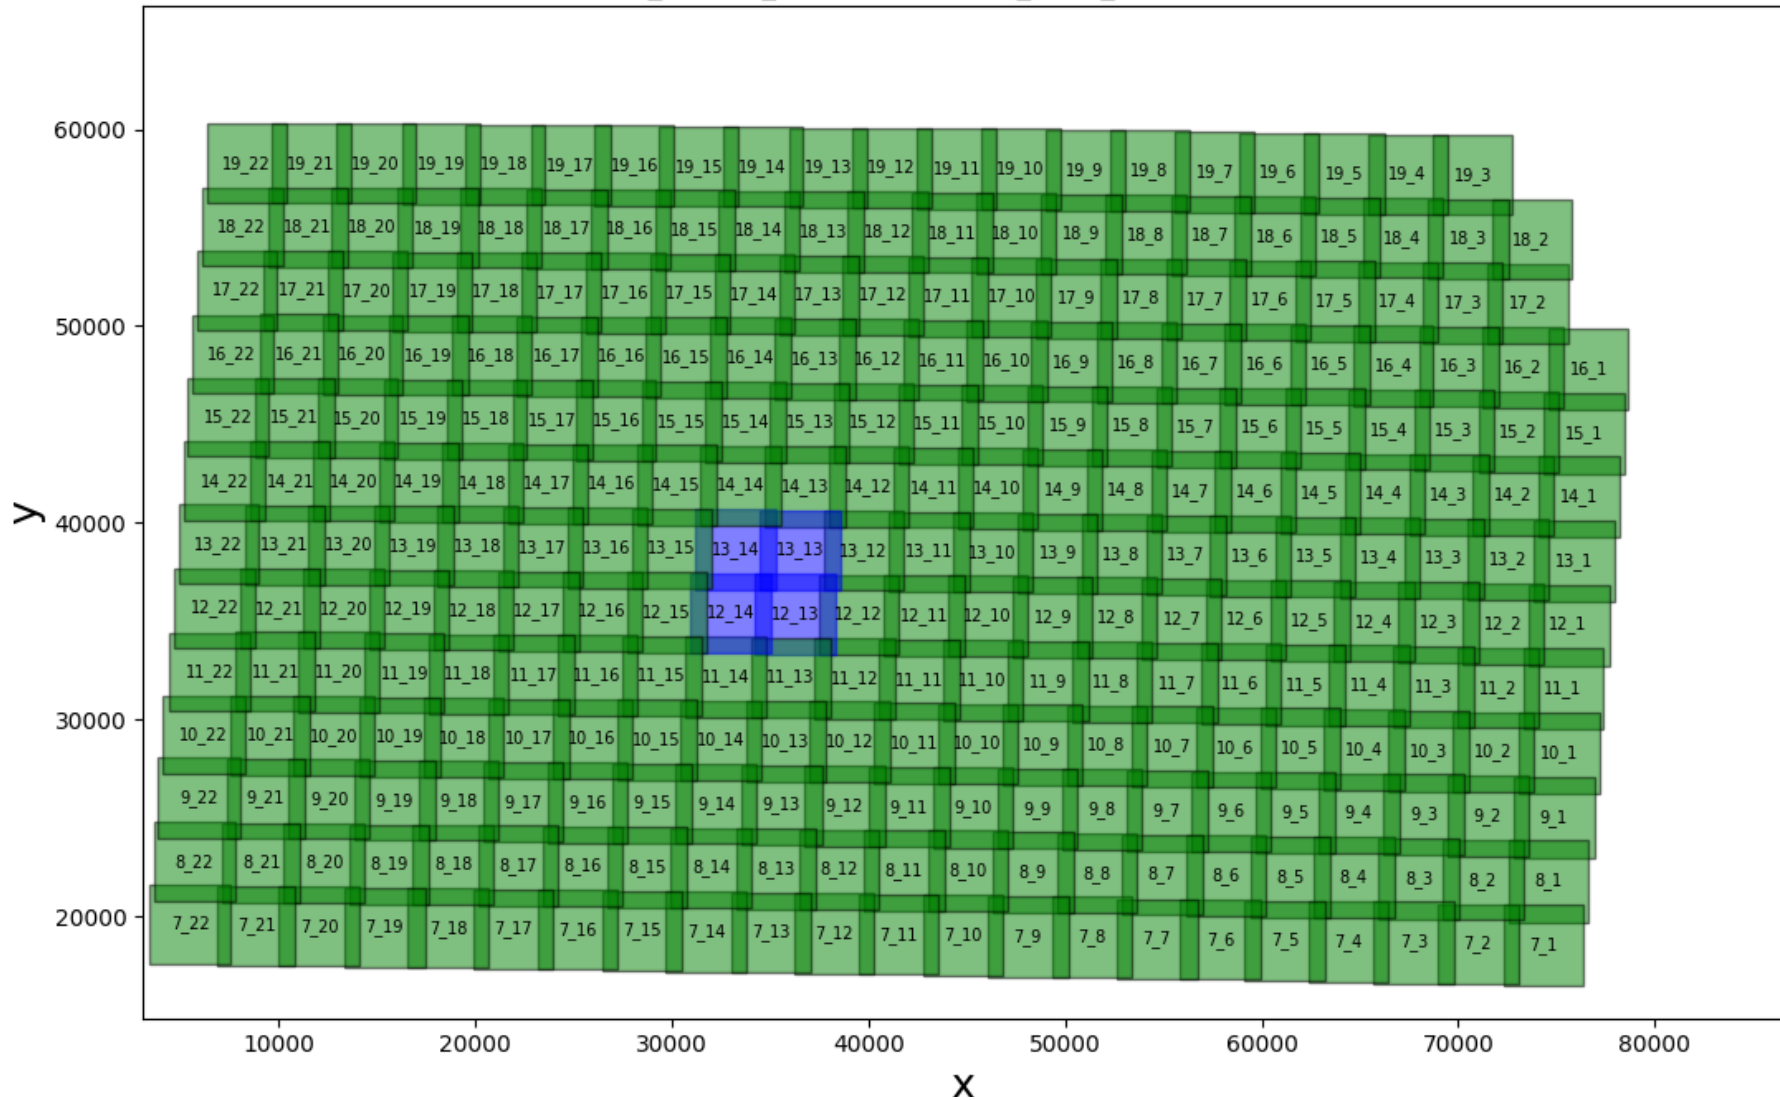
4 Tile example and more

Purpose:

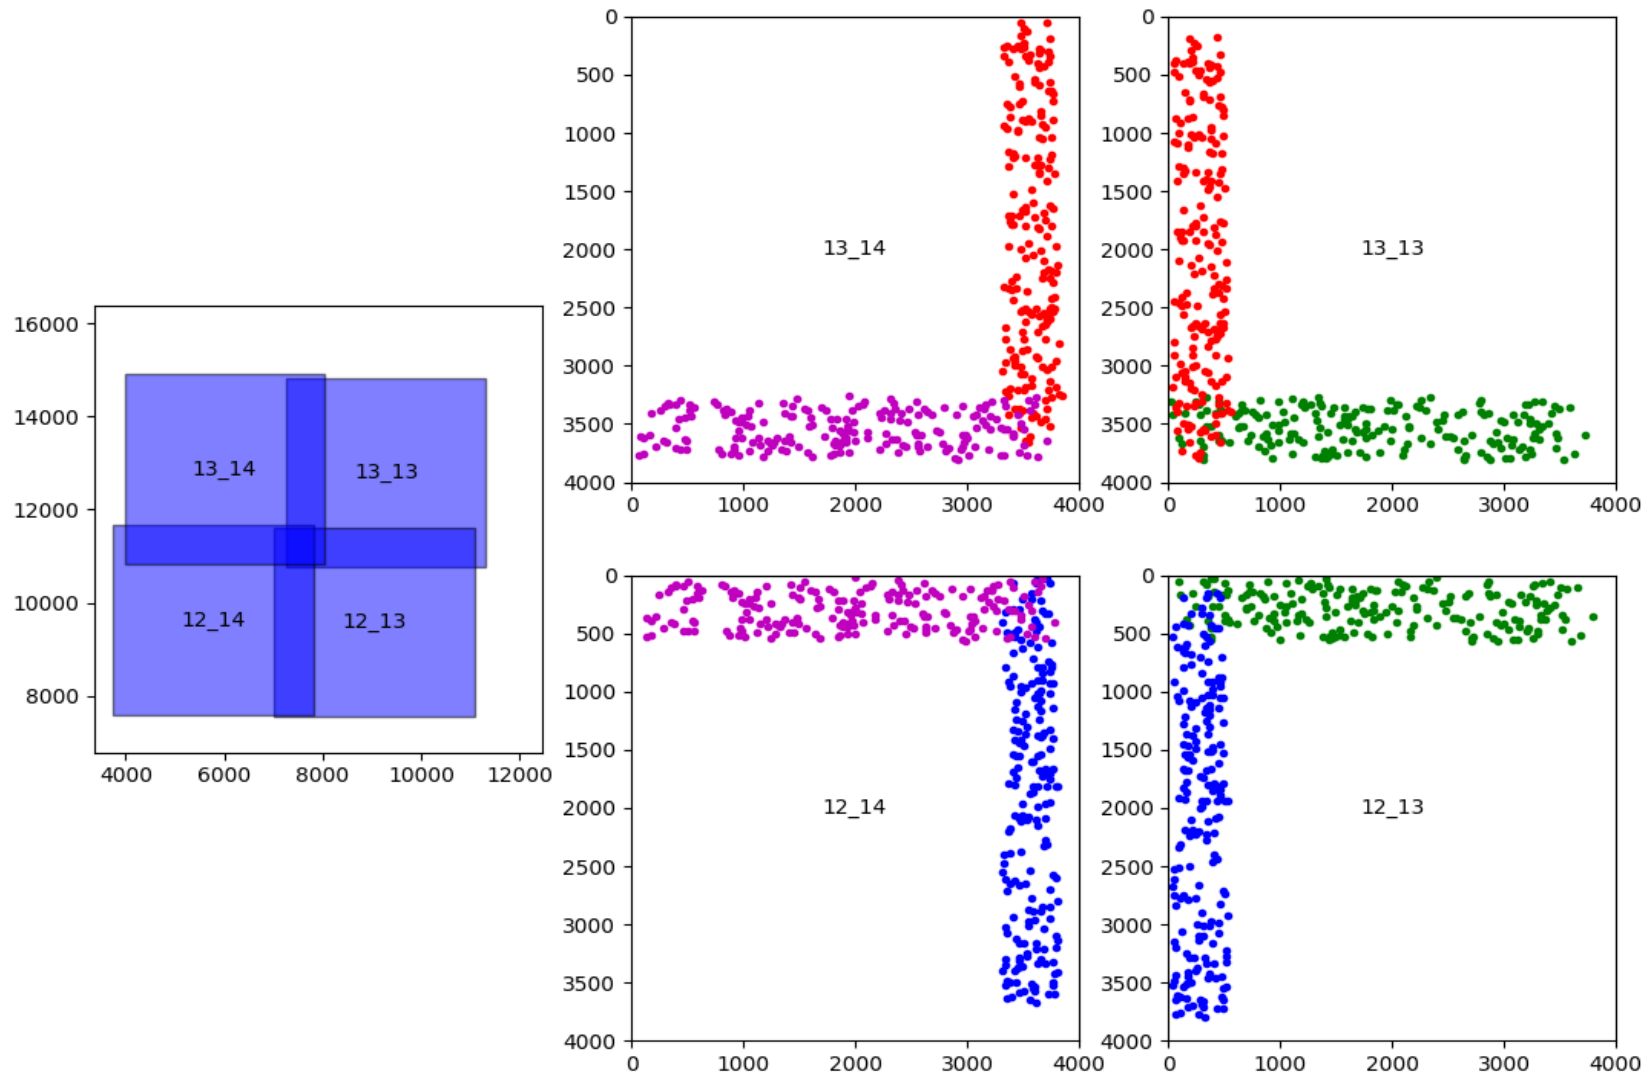
- Understand EM_aligner, MATLAB code base for assembling and solving EM-connectomics fine alignment problems.
- Understand the linear algebra going into these solutions.
- Develop plan for scaling the solver for large chunks of mm² sections.

Source of Data

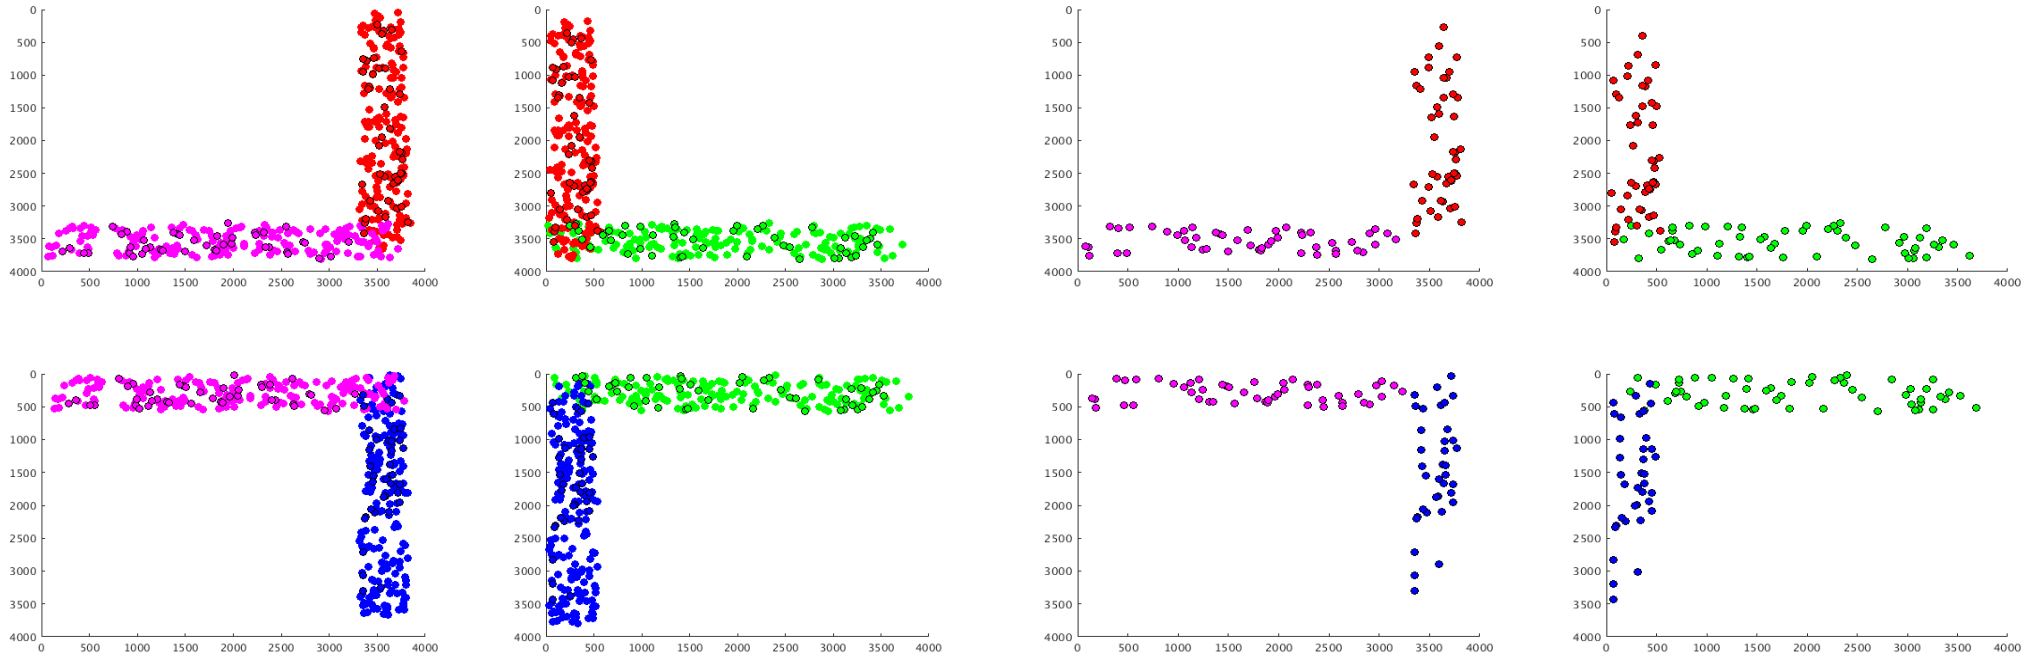
EM_Phase1_Fine | RoughAlign_2266_3483 | z=2316



4 Tiles and Point Matches

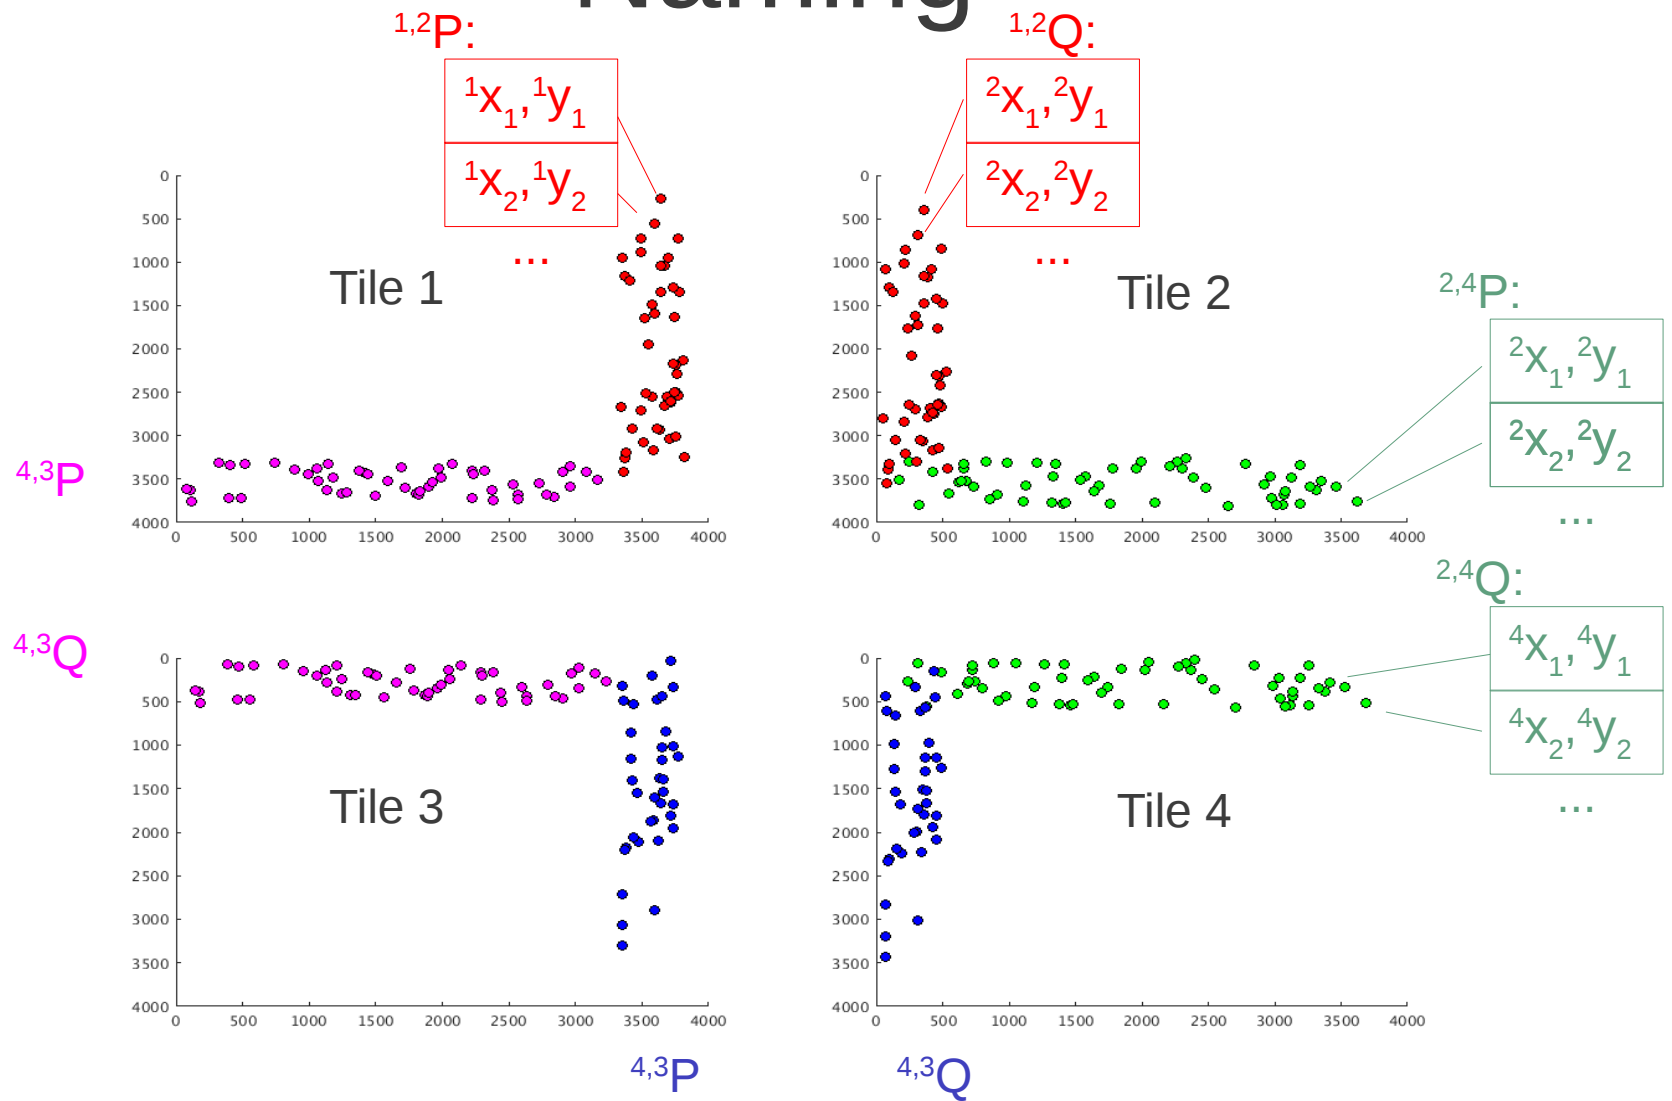


Filtering



- EM_aligner includes filtering to limit the number of included point matches.
- In this example, the filtering reduced the # from 200 points per tile pair to ~50.

Naming



Affine Transformation

The point match coordinates in each tile will be transformed according to an affine transformation:

$$\begin{aligned} {}^1u_1 &= a_1 {}^1x_1 + b_1 {}^1y_1 + c_1 \\ {}^1v_1 &= d_1 {}^1x_1 + e_1 {}^1y_1 + f_1 \end{aligned}$$

The goal of alignment between two tiles is to meet the criteria:

$$\begin{aligned} {}^1u_1 &= {}^2u_1 \\ {}^1v_1 &= {}^2v_1 \\ {}^1u_2 &= {}^2u_2 \\ {}^1v_2 &= {}^2v_2 \end{aligned}$$

...

Explicitly:

$$\begin{aligned} a_1 {}^1x_1 + b_1 {}^1y_1 + c_1 &= a_2 {}^2x_1 + b_2 {}^2y_1 + c_2 \\ d_1 {}^1x_1 + e_1 {}^1y_1 + f_1 &= d_2 {}^2x_1 + e_2 {}^2y_1 + f_2 \end{aligned}$$

Matrix Assembly

$$a_1^1 x_1 + b_1^1 y_1 + c_1 = a_2^2 x_1 + b_2^2 y_1 + c_2$$

$$d_1^1 x_1 + e_1^1 y_1 + f_1 = d_2^2 x_1 + e_2^2 y_1 + f_2$$

which is the same as:

$$\begin{bmatrix} -^1x_1 & -^1y_1 & -1 & ^2x_1 & ^2y_1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & -^1x_1 & -^1y_1 & -1 & ^2x_1 & ^2y_1 & 1 \end{bmatrix} \times \begin{bmatrix} a_1 \\ b_1 \\ c_1 \\ a_2 \\ b_2 \\ c_2 \\ d_1 \\ e_1 \\ f_1 \\ d_2 \\ e_2 \\ f_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

And for multiple pairs of points:

$$\begin{bmatrix} -^1x_1 & -^1y_1 & -1 & ^2x_1 & ^2y_1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -^1x_2 & -^1y_2 & -1 & ^2x_2 & ^2y_2 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -^1x_3 & -^1y_3 & -1 & ^2x_3 & ^2y_3 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -^1x_4 & -^1y_4 & -1 & ^2x_4 & ^2y_4 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & 0 & 0 & 0 & -^1x_1 & -^1y_1 & -1 & ^2x_1 & ^2y_1 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & -^1x_2 & -^1y_2 & -1 & ^2x_2 & ^2y_2 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & -^1x_3 & -^1y_3 & -1 & ^2x_3 & ^2y_3 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & -^1x_4 & -^1y_4 & -1 & ^2x_4 & ^2y_4 & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots & \dots \end{bmatrix} \times \begin{bmatrix} a_1 \\ b_1 \\ c_1 \\ a_2 \\ b_2 \\ c_2 \\ d_1 \\ e_1 \\ f_1 \\ d_2 \\ e_2 \\ f_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ \dots \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ \dots \end{bmatrix}$$

Reorder to keep all tile 1
coordinates in a column block

Matrix Assembly

$$\begin{bmatrix}
 -^1x_1 & -^1y_1 & -1 & 0 & 0 & 0 & ^2x_1 & ^2y_1 & 1 & 0 & 0 & 0 \\
 -^1x_2 & -^1y_2 & -1 & 0 & 0 & 0 & 0 & ^2x_2 & ^2y_2 & 1 & 0 & 0 \\
 -^1x_3 & -^1y_3 & -1 & 0 & 0 & 0 & 0 & ^2x_3 & ^2y_3 & 1 & 0 & 0 \\
 -^1x_4 & -^1y_4 & -1 & 0 & 0 & 0 & 0 & ^2x_4 & ^2y_4 & 1 & 0 & 0 \\
 0 & 0 & 0 & \dots & -^1x_1 & -^1y_1 & -1 & 0 & 0 & 0 & ^2x_1 & ^2y_1 & 1 \\
 0 & 0 & 0 & -^1x_2 & -^1y_2 & -1 & 0 & 0 & 0 & 0 & ^2x_2 & ^2y_2 & 1 \\
 0 & 0 & 0 & -^1x_3 & -^1y_3 & -1 & 0 & 0 & 0 & 0 & ^2x_3 & ^2y_3 & 1 \\
 0 & 0 & 0 & -^1x_4 & -^1y_4 & -1 & 0 & 0 & 0 & 0 & ^2x_4 & ^2y_4 & 1 \\
 & & & & & & & \dots & & & & &
 \end{bmatrix}
 \times
 \begin{bmatrix}
 a_1 \\ b_1 \\ c_1 \\ d_1 \\ e_1 \\ f_1 \\ a_2 \\ b_2 \\ c_2 \\ d_2 \\ e_2 \\ f_2
 \end{bmatrix}
 =
 \begin{bmatrix}
 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0
 \end{bmatrix}$$

Block-wise shorthand:

$$\begin{bmatrix} {}^{1,2}P & -{}^{1,2}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

2 tiles:

$$\begin{bmatrix} {}^{1,2}P & -{}^{1,2}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

3 tiles:

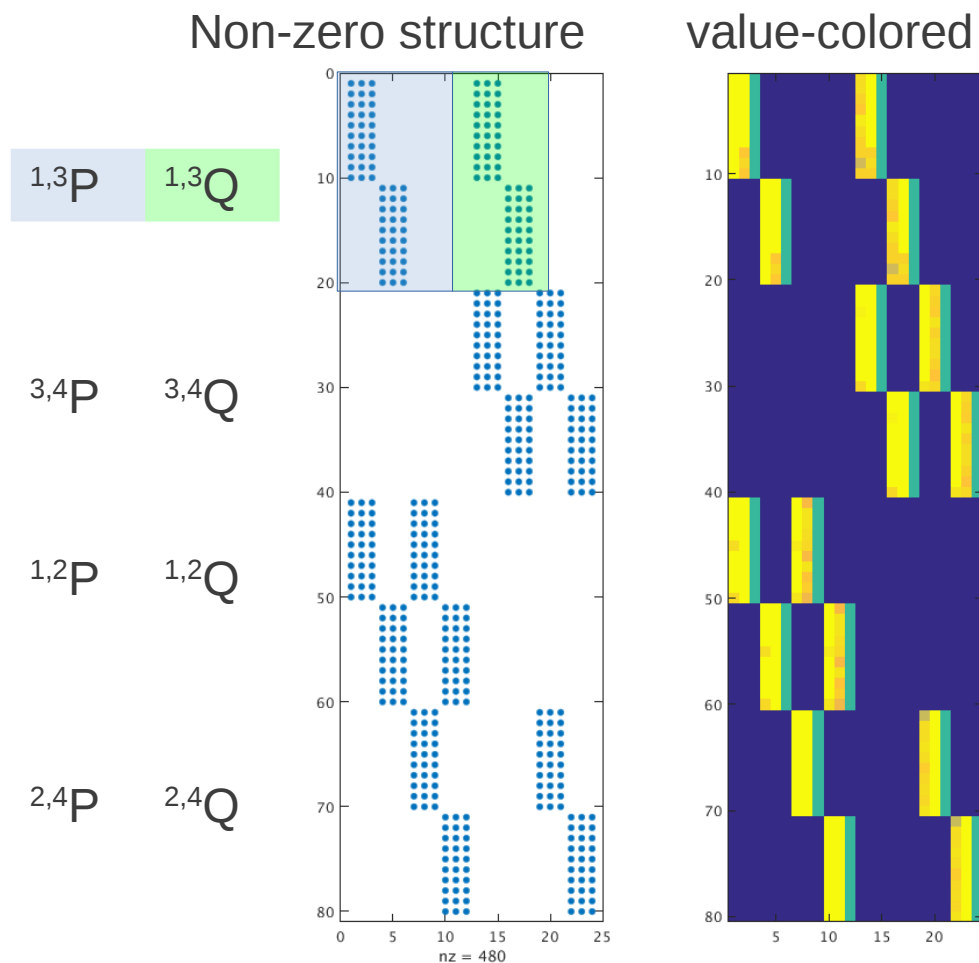
$$\begin{bmatrix} {}^{1,2}P & -{}^{1,2}Q & 0 \\ {}^{1,3}P & 0 & -{}^{1,3}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \\ T_3 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

4 tiles:

$$\begin{bmatrix} {}^{1,2}P & -{}^{1,2}Q & 0 & 0 \\ {}^{1,3}P & 0 & -{}^{1,3}Q & 0 \\ 0 & ^{2,4}P & 0 & -{}^{2,4}Q \\ 0 & 0 & ^{3,4}P & -{}^{3,4}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \\ T_3 \\ T_4 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Matrix Assembly (from data)

Force filtering to limit to 10 point matches per tile, for easy visualization:



The ordering is a little different than the example, but, it's just block-wise row permutations:

$$\begin{bmatrix} {}^{1,3}P & 0 & -{}^{1,3}Q & 0 \\ 0 & 0 & {}^{3,4}P & -{}^{3,4}Q \\ {}^{1,2}P & -{}^{1,2}Q & 0 & 0 \\ 0 & {}^{2,4}P & 0 & -{}^{2,4}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \\ T_3 \\ T_4 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

($Ax = b$)



We know these values. They are the coordinates of the point matches.



We want to know these transform values, and apply them to the image data.

Least Squares

For this simple problem, there are:

- 80 equations (4 tile matches x 10 points x 2 DOF)
- 24 unknowns (4 tiles x (2 affine +1 translation) x 2 DOF)

The system is overdetermined.

Use linear least-squares to determine the best possible set of transformations.

$$\begin{bmatrix} {}^{1,3}P & 0 & -{}^{1,3}Q & 0 \\ 0 & 0 & {}^{3,4}P & -{}^{3,4}Q \\ {}^{1,2}P & -{}^{1,2}Q & 0 & 0 \\ 0 & {}^{2,4}P & 0 & -{}^{2,4}Q \end{bmatrix} \times \begin{bmatrix} T_1 \\ T_2 \\ T_3 \\ T_4 \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

($Ax = b$)

Per row (i) residual: $r_i = b_i - \sum_{j=1}^n A_{ij}x_j$

We want norm² to be as small as possible: $S = \sum_{i=1}^m r_i^2$

Least Squares

Per row (i) residual: $r_i = b_i - \sum_{j=1}^n A_{ij}x_j$

We want norm² to be as small as possible: $S = \sum_{i=1}^m r_i^2$

Find the minimum by seeking where the gradient relative to the parameters = 0

gradient $\frac{\partial S}{\partial x_j} = 2 \sum_i^m r_i \frac{\partial r_i}{\partial x_j}$

$$\frac{\partial r_i}{\partial x_j} = -A_{ij}$$

gradient,
expanded

$$\frac{\partial S}{\partial x_j} = 2 \sum_i^m \left(b_i - \sum_{k=1}^n A_{ik}x_k \right) (-A_{ij}) = 0$$

rearranged

$$\sum_{i=1}^m \sum_{k=1}^n A_{ij}A_{ik}x_k = \sum_{i=1}^m A_{ij}b_i$$

Matrix notation

$$\underline{A^T A} x = A^T b$$

This product is a square matrix (24x24, in our example).

Now there are 24 equations and 24 unknowns.

Weighting and Regularization

Our new matrix equation:

$$A^T A x = A^T b$$

We can weight each row of A differently

$$A^T W A x = A^T W b$$

Example: within-section point matches should be weighted more than cross-section point matches

We can regularize

$$A^T W A x + \lambda = A^T W b + \lambda d$$


Derivation to-be-added, but, it's much like on the last page, with an additional term in S that increases S if the solution strays far from a first guess/approximation for x

Renaming things:


$$K x = L_m$$

Properties of K and L_m

K is symmetric




The graph of K shows two separate sub-matrices.




10/5/2017

Dan Kapner

L_m has non-zero values



Which we can see easily from re-ordering rows and columns:



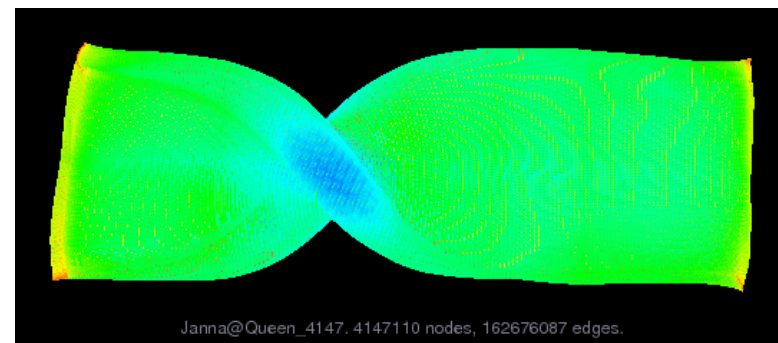

This type of structure will be important for partitioning large matrices for distributed solving. I know about 2 partitioning packages for this:

- Scotch: used by the Pastix solver. (Also, PT-Scotch for parallel partitioning.)
- METIS. Also, (ParMetis)

13

Size estimate for 1mm³

- Estimate for K
 - Ntiles = $2.5e8 = (1e4 \text{ tiles/section}) \times (2.5e4 \text{ sections})$
 - $K \text{ m} \times \text{n} = 1.5e9 \times 1.5e9$ (6DOF per tile)
 - Guess that each tile has point matches to 10 other tiles.
 - nnz = $5e10$ (each row has $3 \times (10+1)$ nz entries)
- Similar (?), smaller matrix: Janna/Queen_4147 (SuiteSparse...):
 - $\text{m} \times \text{n} = 4.1e6 \times 4.1e6$
 - nnz = $3.2e8$



First pastix with Janna/Queen_4147

```
root@15884987838d:/app/pastix_5.2.3/src/example/bin# ./simple -MM /app/matrices/Queen_4147.mtx
```

```
MPI_Init_thread level = MPI_THREAD_MULTIPLE
```

```
driver: MatrixMarket file: /app/matrices/Queen_4147.mtx
```

```
open RHS file : /app/matrices/Queen_4147.mtx.rhs
```

```
cannot load /app/matrices/Queen_4147.mtx.rhs
```

```
Setting right-hand-side member such as X[i] = i
```

```
Check : Numbering          OK
```

```
Check : Sort CSC           OK
```

```
Check : Duplicates         OK
```


```
+-----+
+           PaStiX : Parallel Sparse matrix package           +
+-----+
Matrix size                      4147110 x 4147110
Number of nonzeros in A          166823197
+-----+
+ Options                                                         +
+-----+
Version      :
SMP_SOPALIN  : Defined
VERSION MPI  : Defined
PASTIX_DYNSED : Not defined
STATS_SOPALIN : Not defined
NAPA_SOPALIN : Defined
TEST_IRecv   : Not defined
TEST_ISEND   : Defined
TAG          : Exact Thread
FORCE_CONSO  : Not defined
RECV_FANIN_OR_BLOCK : Not defined
OUT_OF_CORE  : Not defined
DISTRIBUTED  : Defined
METIS        : Not defined
WITH_SCOTCH  : Defined
INTEGER TYPE : int
PASTIX_FLOAT TYPE : double
+-----+
Time to compute ordering          46.2 s
Time to analyze                   5.09 s
Number of nonzeros in factorized matrix -2147483648
Fill-in                          -12.8728
Number of operations (LLt)        2.69604e+14
Prediction Time to factorize (AMD 6180 MKL) 3.61e+04 s
--- Sopalin : Threads are binded ---
```

First, try simple.c (single-threaded...)


- 1.7e8 nnz before factorization
- -2.1e9 nnz after factorization
(note minus sign, I thought I compiled with 64-bit integers, but, now I need to go check)
- At start of factorization, RAM usage went up to 115GB. (from ~5GB)
- 10 hours predicted time to factorize...
- This is a terrible example for single-thread, but, good for parallel

A and K from a whole section (282 tiles)


structure of A




structure of K




graph of K



Do these plots make sense?



Change aspect



A size: 50532x1692
1692 = 282tiles x 6DOF/tile

50532 = average of ~190
point matches per overlap